

Assembly Crash Course

System Calls

Yan Shoshitaishvili
Arizona State University

Having Effects

```
exit();
```

How do we interact with the outside world?

Even something as simple as quitting the program?

System Calls

Remember system calls? It's an instruction that makes a *call* into the Operating System.

`syscall` triggers the system call specified by the value in `rax`.

arguments in `rdi`, `rsi`, `rdx`, `r10`, `r8`, and `r9`

return value in `rax`

Reading 100 bytes from stdin to the stack:

```
n = read(0, buf, 100);
```

```
mov rdi, 0 # the stdin file descriptor
```

```
mov rsi, rsp # read the data onto the stack
```

```
mov rdx, 100 # the number of bytes to read
```

```
mov rax, 0 # system call number of read()
```

```
syscall # do the system call
```

`read` returns the number of bytes read via `rax`, so we can easily **write** them out:

```
write(1, buf, n);
```

```
mov rdi, 1 # the stdout file descriptor
```

```
mov rsi, rsp # write the data from the stack
```

```
mov rdx, rax # the number of bytes to write (same as what we read in)
```

```
mov rax, 1 # system call number of write()
```

```
syscall # do the system call
```

System Calls

System calls have very well-defined interfaces that very rarely change.

There are over 300 system calls in Linux. Here are some examples:

`int open(const char *pathname, int flags)` - returns a file new file descriptor of the open file (also shows up in `/proc/self/fd!`)

`ssize_t read(int fd, void *buf, size_t count)` - reads data from the file descriptor

`ssize_t write(int fd, void *buf, size_t count)` - writes data to the file descriptor

`pid_t fork()` - forks off an *identical* child process. Returns 0 if you're the child and the PID of the child if you're the parent.

`int execve(const char *filename, char **argv, char **envp)` - *replaces* your process.

`pid_t wait(int *wstatus)` - wait child termination, return its PID, write its status into `*wstatus`.

Look familiar?

"String" Arguments

Some system calls take "string" arguments (for example, file paths).

A string is a bunch of contiguous bytes in memory, followed by a `0` byte.

Let's build a file path for `open()` on the stack:

```
mov BYTE PTR [rsp+0], '/' # write the ASCII value of / onto the stack
```

```
mov BYTE PTR [rsp+1], 'f'
```

```
mov BYTE PTR [rsp+2], 'l'
```

```
mov BYTE PTR [rsp+3], 'a'
```

```
mov BYTE PTR [rsp+4], 'g'
```

```
mov BYTE PTR [rsp+5], 0 # write the 0 byte that will terminate our string
```

rsp	rsp+1	rsp+2	rsp+3	rsp+4	rsp+5
2f (/)	66 (f)	6c (l)	61 (a)	67 (g)	00 (\0)

Now, we can `open()` the `/flag` file!

```
mov rdi, rsp # read the data onto the stack
```

```
mov rsi, 0 # open the file read only (more on this later)
```

```
mov rax, 2 # system call number of open()
```

```
syscall # do the system call
```

`open()` returns the file descriptor number in `rax`

Constant Arguments

The argument `flags` must include one of the following access modes: `O_RDONLY`, `O_WRONLY`, or `O_RDWR`. These request opening the file read-only, write-only, or read/write, respectively.

Some system calls require archaic "constants".

Example: `open()` has a `flags` argument to determine how the file will be opened.

We can figure out the values of these arguments in C!

```
#include <stdio.h>
#include <fcntl.h>
int main() {
    printf("O_RDONLY is: %d\n", O_RDONLY);
}
```

```
yans@ramoth ~/pwn $ ./print_rdonly
O_RDONLY is: 0
yans@ramoth ~/pwn $
```

Quitting The Program

Finally, we can quit!

```
mov rdi, 42 # our program's return code (e.g., for bash scripts)
mov rax, 60 # system call number of exit()
syscall # do the system call
```

Goodbye, world!